
CPAEP AI Summary of Papers

An Open-Source Summary

Paper's author and Gemini



February 28, 2026

Contents

1	L1 - The Apple M1 Architecture	3
1.1	Performance and Efficiency Analysis	3
1.2	Strategic Implications	3
2	L1 - Summary: A New Golden Age for Computer Architecture	3
2.1	1. Historical Context: From CISC to RISC	4
2.2	2. Current Challenges	4
2.3	3. Future Opportunities	4
	2.3.1 Domain-Specific Architectures (DSAs)	4
	2.3.2 Open Instruction Sets	4
	2.3.3 Agile Hardware Development	5
2.4	Conclusion	5
3	L2 - Summary: Attention Is All You Need	5
3.1	Overview	5
3.2	Key Innovations	5
	3.2.1 1. The Transformer Architecture	5
	3.2.2 2. Multi-Head Attention	5
	3.2.3 3. Positional Encoding	6
3.3	Why Self-Attention?	6
3.4	Performance and Results	6
3.5	Conclusion	6
4	L3 - Paper Summary: “How to Keep Pushing ML Accelerator Performance? Know Your Rooflines!”	6
	4.0.1 1. The Dual Roofline Framework	7
	4.0.2 2. Key Efficiency-Boosting Techniques	7
	4.0.3 3. Critical Trade-offs: The “Balancing Game”	7
	4.0.4 4. Conclusion and Future Outlook	8
5	L4 - Executive Summary: In-Datcenter Performance Analysis of a Tensor Processing Unit	8
	5.0.1 Key Architectural Features	8
	5.0.2 Performance and Efficiency Findings	8
	5.0.3 Critical Insights	9
	5.0.4 Conclusion	9
6	L5 - Summary: ENVISION - A High-Efficiency CNN Processor	9
6.1	Overview	9
6.2	Key Innovations	9
	6.2.1 1. Hierarchical Recognition Processing	9
	6.2.2 2. Subword-Parallel DVAFS	10
	6.2.3 3. Body-Bias (BB) Modulation	10
6.3	Technical Specifications & Architecture	10
6.4	Results and Comparison	10

7	L6 - Paper Summary: DIANA - A Hybrid Digital/Analog Neural Network SoC	11
7.0.1	1. Architecture Overview	11
7.0.2	2. Key Technical Innovations	11
7.0.3	3. Performance and Efficiency	11
7.0.4	4. Conclusion	12
8	L7 - Paper Summary: XpulpNN: Accelerating Quantized Neural Networks on RISC-V Processors Through ISA Extensions	12
8.1	Overview	12
8.2	Key Contributions	12
8.2.1	1. ISA Extensions (XpulpNN)	12
8.2.2	2. Microarchitectural Optimization	12
8.3	Experimental Results	13
8.3.1	Performance & Efficiency Gains	13
8.3.2	Comparison with State-of-the-Art	13
8.4	Conclusion	13
9	L7 - Summary of ARM big.LITTLE Technology White Paper	13
9.0.1	Core Concept: Heterogeneous Coherency	13
9.0.2	Software Execution Models	14
9.0.3	Dynamic Task Management in big.LITTLE MP	14
9.0.4	Performance and Efficiency Results	14
10	L8 - Summary: A RISC-V Vector Processor with Tightly-Integrated Switched-Capacitor DC-DC Converters	15
10.1	Overview	15
10.2	Key Technical Features	15
10.2.1	1. RISC-V Processor Architecture	15
10.2.2	2. Integrated Power Management	15
10.2.3	3. Adaptive Clocking System	15
10.3	Measured Performance and Results	16
10.4	Conclusion	16
11	L9 - Summary: elQ Neutron - Redefining Edge-AI Inference	16
11.1	Overview	16
11.2	Key Architectural Principles	16
11.3	Innovations	17
11.3.1	1. Hardware Architecture	17
11.3.2	2. Compiler and Software Stack	17
11.4	Experimental Results	17
11.5	Conclusion	17
12	License	18
12.1	Modification	18
12.2	Take down	18

1 L1 - The Apple M1 Architecture

The M1 is Apple's first System on a Chip (SoC) designed specifically for Macs, utilizing a **5nm process node** with **16 billion transistors**. Key architectural features include:

- **CPU Configuration:** An 8-core design consisting of **4 “Firestorm” performance cores** and **4 “Icestorm” efficiency cores**.
 - **Unified Memory Architecture (UMA):** Uses an Apple-designed package with high-bandwidth, low-latency DRAM accessible to the entire SoC.
 - **Execution Width:** The Firestorm cores feature an industry-leading **8-wide decode block**, significantly wider than contemporary designs from Intel and AMD.
 - **Cache Hierarchy:** Includes a massive **192KB L1 instruction cache** and a shared **12MB L2 cache** for the performance cores—a 50% increase over the A14.
-

1.1 Performance and Efficiency Analysis

The report highlights a dramatic shift in the competitive landscape between Apple and traditional x86 vendors:

- **Performance Benchmarks:** Testing shows the A14/M1 architecture competes with, and often beats, the best desktop and laptop CPUs from Intel and AMD in single-threaded performance.
 - **Power Advantage:** Apple achieves this performance at a fraction of the power consumption. For example, the M1 is estimated to peak at **18W**, while competing PC chips peak between **35-40W**.
 - **Trajectory:** Over the last five years, Apple has improved its single-thread performance by nearly **300% (2.98x)**, while Intel managed only a **28%** increase in the same period.
-

1.2 Strategic Implications

The author concludes that Apple's “vertical integration” across hardware and software allows for a swift transition that other companies cannot easily replicate. By ditching Intel, Apple avoids stagnation and gains the ability to introduce new technologies (like integrated AI and security controllers) without being held back by third-party chip roadmaps.

2 L1 - Summary: A New Golden Age for Computer Architecture

In their Turing Lecture, John L. Hennessy and David A. Patterson review the history of computer architecture and argue that the field is entering a “New Golden Age” driven by the end of traditional scaling laws and the rise of domain-specific solutions.

2.1 1. Historical Context: From CISC to RISC

- **The Unification Era:** In the 1960s, IBM's System/360 pioneered the use of a single Instruction Set Architecture (ISA) to unify diverse computer lines, utilizing microprogramming to simplify complex control hardware.
- **The CISC Era:** Integrated circuits enabled larger control stores, leading to Complex Instruction Set Computers (CISC) like the VAX and Intel 8086. However, the marketplace often favored “emergency” designs (like the 8086) over more ambitious ones (like the iAPX-432).
- **The RISC Revolution:** Research at Berkeley and Stanford in the 1980s demonstrated that simpler, Reduced Instruction Set Computers (RISC) could be significantly faster by eliminating microcode interpreters and focusing on compiler-driven register allocation.

2.2 2. Current Challenges

The paper identifies several “insoluble problems” that are now facts of life for modern architects:

- **End of Moore's Law:** Transistor density is no longer doubling at its historical rate, creating a widening gap between predicted and actual capability.
- **End of Dennard Scaling:** Power consumption per transistor is no longer dropping as density increases, leading to the “dark silicon” era where thermal limits prevent all parts of a chip from running simultaneously.
- **Amdahl's Law:** Performance gains from multicore processors are strictly limited by the sequential portions of a program.
- **Security Flaws:** Modern performance techniques like speculative execution have introduced massive vulnerabilities (e.g., Spectre and Meltdown) that exploit the hardware implementation rather than software bugs.

2.3 3. Future Opportunities

Hennessy and Patterson project that the next decade will see a “Cambrian explosion” of architecture innovation through three primary pillars:

2.3.1 Domain-Specific Architectures (DSAs)

Rather than general-purpose CPUs, architects are moving toward processors tailored to specific domains (like the Google TPU for machine learning). DSAs achieve efficiency by:

- Exploiting more efficient parallelism (SIMD).
- Using software-controlled local memories instead of complex, energy-hungry caches.
- Reducing data precision (e.g., using 8-bit integers) when full 64-bit accuracy is unnecessary.

2.3.2 Open Instruction Sets

The rise of open ISAs, most notably **RISC-V**, allows for a “Linux for processors.” This modular, simple, and clean-slate design enables:

- Public collaboration and faster innovation.
- The ability to add custom extensions for accelerators without proprietary restrictions.
- Better security through transparency and simplicity.

2.3.3 Agile Hardware Development

Inspired by software development, the paper advocates for agile hardware methodologies. By utilizing:

- Higher levels of abstraction in electronic CAD tools.
- Cloud-based FPGAs for rapid prototyping.
- Inexpensive small-chip fabrication (e.g., 28nm “tape-ins”). Architects can iterate on designs in weeks rather than years, accelerating the path from research to commercial adoption.

2.4 Conclusion

The authors conclude that the convergence of domain-specific languages, open architectures, and agile development will lead to a new era of gains in cost, energy, and security, marking a second golden age for the field of computer architecture.

3 L2 - Summary: Attention Is All You Need

3.1 Overview

“Attention Is All You Need” introduces the **Transformer**, a novel network architecture for sequence transduction that relies entirely on **self-attention mechanisms**. Unlike previous dominant models that utilized complex recurrent neural networks (RNNs) or convolutional neural networks (CNNs), the Transformer dispenses with recurrence and convolutions entirely. This shift allows for significantly more parallelization during training and has established new state-of-the-art benchmarks in translation quality.

3.2 Key Innovations

3.2.1 1. The Transformer Architecture

The model follows a standard **encoder-decoder** structure:

- **Encoder:** Composed of a stack of 6 identical layers. Each layer contains two sub-layers: a multi-head self-attention mechanism and a position-wise fully connected feed-forward network.
- **Decoder:** Also composed of 6 identical layers. In addition to the two sub-layers found in the encoder, the decoder inserts a third sub-layer to perform multi-head attention over the encoder’s output.
- **Residual Connections:** Each sub-layer utilizes residual connections followed by layer normalization.

3.2.2 2. Multi-Head Attention

Instead of performing a single attention function, the Transformer uses **Multi-Head Attention**. This allows the model to jointly attend to information from different representation subspaces at different positions. It consists of:

- **Scaled Dot-Product Attention:** The core mechanism that computes attention weights by scaling the dot product of queries and keys.
- **Parallel Heads:** The model employs 8 parallel attention heads to learn diverse dependencies.

3.2.3 3. Positional Encoding

Since the model lacks recurrence or convolution, it has no inherent sense of the order of the sequence. To address this, **positional encodings** (using sine and cosine functions) are added to the input embeddings to provide information about the relative or absolute position of tokens.

3.3 Why Self-Attention?

The authors motivate the use of self-attention based on three primary factors:

1. **Computational Complexity:** Self-attention layers are generally faster than recurrent layers when the sequence length is shorter than the representation dimensionality.
 2. **Parallelization:** It minimizes the number of sequential operations required, significantly reducing training time.
 3. **Long-Range Dependencies:** Self-attention reduces the path length between any two positions to a constant $O(1)$ operations, making it easier for the model to learn dependencies regardless of their distance in the sequence.
-

3.4 Performance and Results

The Transformer was evaluated on two major machine translation tasks:

- **WMT 2014 English-to-German:** The “big” model achieved a state-of-the-art **28.4 BLEU** score, outperforming previous best results and ensembles by over 2 BLEU.
- **WMT 2014 English-to-French:** The model reached a score of **41.0 BLEU**, establishing a new single-model state-of-the-art at a fraction of the training cost of previous models (3.5 days on 8 GPUs).

3.5 Conclusion

The Transformer represents a major shift in NLP, proving that attention alone is sufficient for high-quality sequence modeling. It offers superior translation quality, better parallelization, and faster training speeds compared to its recurrent and convolutional predecessors.

4 L3 - Paper Summary: “How to Keep Pushing ML Accelerator Performance? Know Your Rooflines!”

This paper, published in the *Journal of Solid State Circuits* (2024), explores the rapid evolution of Machine Learning (ML) hardware accelerators. The authors provide a unifying framework to understand how architectural innovations—specifically

parallelism, data reuse, quantization, sparsity, and in-memory computing—interact to overcome the limitations of Moore’s Law.

4.0.1 1. The Dual Roofline Framework

The authors introduce an enhanced version of the **Roofline Model** to visualize design trade-offs. Unlike traditional models that focus only on throughput, this paper emphasizes two distinct boundaries:

- **Performance Roofline:** Charts attainable operations per cycle against Arithmetic Intensity (AI). It identifies if a system is **memory-bound** (limited by bandwidth) or **compute-bound** (limited by the number of multipliers).
- **Energy Roofline:** Charts energy efficiency (TOPS/W). Because energy is additive rather than a “max/min” function, this roofline is curved, showing how memory access energy dominates at low AI and compute energy dominates at high AI.

4.0.2 2. Key Efficiency-Boosting Techniques

The paper surveys five primary strategies used by state-of-the-art (SOTA) accelerators like Google’s TPU and NVIDIA’s GPUs:

- **Maximizing Parallelism:** Increasing the number of Multiply-Accumulate (MAC) units. While this raises the “roof,” it shifts the knee point to the right, making the system more prone to being memory-bound.
 - **Data Reuse:** Leveraging spatial and temporal reuse to increase AI. By keeping weights or activations stationary in local registers, the hardware performs more compute per memory fetch.
 - **Quantization:** Reducing numerical precision (e.g., from FP32 to INT8 or MXFP4). This simultaneously raises the compute peak and increases AI by packing more operands into each memory word.
 - **Sparsity:** Skipping operations involving zero-valued weights or activations. While effective for reducing total operations, it often leads to “detachment” from the roofline due to the overhead of indexing and irregular memory access.
 - **Near- and In-Memory Computing (IMC):** Fusing storage and logic to eliminate the energy cost of moving data across a chip. The paper highlights **Switched-Capacitor (SC)** IMC as a high-SNR approach to achieve extreme efficiency (up to 120 TOPS/W).
-

4.0.3 3. Critical Trade-offs: The “Balancing Game”

The authors argue that modern design is no longer just about raising the roofline, but about **operating close to it**. They identify several conflicting goals:

- **Parallelism vs. Utilization:** Massive compute arrays (like 256x256 grids) offer high peak performance but suffer from low “spatial utilization” if the workload dimensions don’t match the hardware.
- **Flexibility vs. Specialization:** Ultra-specialized accelerators (NPUs) achieve peak efficiency for specific kernels (like CNNs) but may become obsolete or under-utilized as models evolve toward Transformers or newer topologies.
- **Storage vs. Compute:** In IMC, coupling memory and logic can lead to storage waste if weights must be replicated to keep compute units busy.

4.0.4 4. Conclusion and Future Outlook

The paper concludes that while hardware and algorithmic improvements have delivered a **1,000x gain** in efficiency over the last eight years, the field is hitting new limits. Future progress will rely on:

- 1. **Chiplet Technology:** For scaling bandwidth and compute parallelism.
- 2. **Advanced Abstractions:** Utilizing frameworks like MLIR to bridge the gap between high-level algorithms and hardware-specific micro-coding.
- 3. **Adaptive Architectures:** Designing hardware that remains “future-proof” by balancing extreme specialization with enough programmability to handle evolving ML workloads.

5 L4 - Executive Summary: In-Datcenter Performance Analysis of a Tensor Processing Unit

This paper evaluates Google’s **Tensor Processing Unit (TPU)**, a custom ASIC designed to accelerate neural network (NN) inference in datacenters. Deployed since 2015, the TPU was developed to address the massive computational demands of services like voice search. The study compares the TPU against contemporary Intel Haswell CPUs and Nvidia K80 GPUs using production workloads (MLPs, CNNs, and LSTMs) that represent 95% of Google’s datacenter NN inference demand.

5.0.1 Key Architectural Features

- **Matrix Multiply Unit:** The core of the TPU is a 65,536 8-bit MAC unit providing a peak throughput of **92 TeraOps/second (TOPS)**.
- **Systolic Execution:** To save energy and reduce memory reads/writes, the matrix unit utilizes a systolic array data flow.
- **Software-Managed Memory:** Features a large 28 MiB on-chip memory (Unified Buffer) and 8 GiB of off-chip Weight Memory.
- **Deterministic Model:** Unlike CPUs and GPUs that optimize for average throughput via caches and out-of-order execution, the TPU uses a deterministic execution model to meet strict **99th-percentile response-time** requirements.
- **CISC Architecture:** Instructions are sent from the host CPU over PCIe, following a CISC tradition to simplify hardware and reduce bus traffic.

5.0.2 Performance and Efficiency Findings

The paper utilizes the **Roofline Performance Model** to visualize bottlenecks, revealing that many applications are memory-bandwidth limited on the TPU despite its high peak performance.

Metric	TPU vs. Haswell CPU	TPU vs. Nvidia K80 GPU
Inference Speed	~15X – 30X faster	~15X – 30X faster
Energy Efficiency (TOPS/Watt)	~30X – 80X higher	~30X – 80X higher

Metric	TPU vs. Haswell CPU	TPU vs. Nvidia K80 GPU
Incremental Perf/Watt	Up to 83X higher	Up to 29X higher

5.0.3 Critical Insights

- **Latency vs. Throughput:** Inference applications often emphasize response time over raw throughput. The TPU operates much closer to its peak performance (80%) while meeting a 7ms latency limit, whereas CPUs and GPUs often waste potential throughput to stay within strict response-time bounds.
- **Workload Reality:** While the academic community focuses heavily on Convolutional Neural Networks (CNNs), these represent only **5%** of Google’s datacenter workload. Multi-Layer Perceptrons (MLPs) and Long Short-Term Memory (LSTMs) dominate the remaining demand.
- **Energy Proportionality:** The TPU shows relatively poor energy proportionality (using 88% of peak power at 10% load) due to its rapid development cycle (15 months), which precluded the inclusion of advanced energy-saving features.
- **Future Headroom:** Modeling suggests that using faster memory (like GDDR5) would triple the TPU’s performance and increase its performance/Watt advantage to 70X over the GPU and 200X over the CPU.

5.0.4 Conclusion

The TPU demonstrates that domain-specific architectures can provide order-of-magnitude improvements in performance and efficiency. By prioritizing a massive matrix unit and a deterministic execution model over general-purpose features, the TPU successfully met the scaling challenges of modern machine learning at a fraction of the cost and power of traditional hardware.

6 L5 - Summary: ENVISION - A High-Efficiency CNN Processor

6.1 Overview

“ENVISION” is a low-power Convolutional Neural Network (CNN) processor fabricated in 28nm FDSOI technology . The paper addresses the energy bottleneck of advanced ConvNets, which typically require billions of computations, making them difficult to implement in “always-on” wearable devices . Envision achieves a peak energy efficiency of **10 TOPS/W**, providing a significant leap over contemporary mobile GPUs and accelerators .

6.2 Key Innovations

6.2.1 1. Hierarchical Recognition Processing

The platform utilizes a hierarchy of increasingly complex, individually trained ConvNets.

- For example, in face recognition, the system can constantly scan for faces at a very low average energy cost .

- It scales up to more complex networks only when specific identification is required, such as detecting a device's owner or performing full VGG-16-based recognition .

6.2.2 2. Subword-Parallel DVAFS

The core technical contribution is **Subword-parallel Dynamic-Voltage-Accuracy-Frequency Scaling (DVAFS)**.

- **Precision Scaling:** The hardware reconfigures its 16b multiply-accumulate (MAC) units to compute multiple lower-precision operations in parallel, such as $2 \times 8b$ or $4 \times 4b$ per cycle .
- **Energy Savings:** Unlike previous techniques, DVAFS allows lowering the frequency and voltage of the full system, including memory and control units, which drastically reduces non-compute energy overheads at low precision .

6.2.3 3. Body-Bias (BB) Modulation

Leveraging FDSOI technology, Envision tunes the balance between dynamic and leakage power based on computational precision .

- At **high precision**, the supply voltage is reduced to save dynamic energy while maintaining speed .
- At **low precision**, the supply voltage and body bias are increased to minimize leakage overhead at constant speed .

6.3 Technical Specifications & Architecture

- **Technology:** 28nm UTBB FD-SOI with a 1.87mm^2 total area .
- **Compute Array:** A 16×16 2D-SIMD MAC array that generates up to $N \times 256$ intermediate outputs per cycle .
- **Memory:** $64 \times 2kB$ single-port SRAM macros subdivided into parallel blocks .
- **Data Compression:** Includes a Huffman DMA used for compressing I/O bandwidth by up to 5.8x .
- **Sparsity Guarding:** The chip skip memory fetches and MAC operations for zero-valued data, leading to an additional 1.6x gain in energy efficiency .

6.4 Results and Comparison

The Envision chip demonstrates a **40x energy-precision scalability** while maintaining constant throughput .

- **AlexNet:** Efficiency scales between 0.8 and 3.8 TOPS/W .
- **VGG-16:** Achieves an average efficiency of 2 TOPS/W .
- **Power Range:** Scaling from 16-bit to 4-bit sparse computations at 76 GOPS reduces power from 290mW down to 7.6mW .

7 L6 - Paper Summary: DIANA - A Hybrid Digital/Analog Neural Network SoC

This paper, presented at **ISSCC 2022**, introduces **DIANA**, an end-to-end energy-efficient System-on-Chip (SoC) designed for Deep Neural Network (DNN) inference on edge devices. The architecture addresses the trade-off between the high precision/flexibility of digital accelerators and the extreme energy efficiency of Analog In-Memory Computing (AiMC).

7.0.1 1. Architecture Overview

DIANA is a hybrid system that combines three major components to achieve state-of-the-art performance:

- **Digital NN Core:** A 16x16 MAC array supporting symmetrical 2-to-8b precision for high-precision layers (e.g., first/last layers) and complex dataflows like Fully Connected (FC) layers.
- **Analog In-Memory Compute (AiMC) Core:** A large 1152x512 array optimized for computationally intensive convolutional layers using low precision (ternary weights and 7b activations).
- **RISC-V Host Processor:** Acts as the system controller, managing instruction streams, memory traffic, and offloading tasks to the appropriate core.

7.0.2 2. Key Technical Innovations

- **Hybrid Execution Strategy:** DIANA assigns layers based on their characteristics. High-precision or low-channel-count layers run on the digital core, while massive intermediate convolutional layers are offloaded to the AiMC core.
- **Inter-layer Pipelining & Layer Fusion:** The system can propagate small tiles of data across multiple layers without sending intermediate activations back to L2 memory. This reduces memory capacity requirements by up to **7.2x**.
- **Output Unrolling:** To improve utilization, DIANA can expand up to 4 output pixels across AiMC columns when the number of output channels is low, accelerating compute by up to **6.6x** in specific blocks.
- **Shared Memory Subsystem:** An optimized L1 activation memory (256 KB) is shared between cores to facilitate efficient data exchange and layer-fused execution.

7.0.3 3. Performance and Efficiency

Implemented in **22nm FDSOI** technology, DIANA demonstrates significant efficiency gains:

- **Peak Efficiency:** The digital core achieves up to **4.1 TOPS/W**, while the AiMC core alone reaches up to **600 TOPS/W**.
 - **End-to-End Hybrid Performance:** For a ResNet-20 workload, the hybrid solution (digital for first/last layers, AiMC for intermediate) achieves **14.4 TOPS/W**, compared to 2.1 TOPS/W for a purely digital execution.
 - **End-to-End Accuracy:** The chip maintains high accuracy despite analog noise, achieving **89% on CIFAR-10 (ResNet-20)** and **64% on ImageNet (ResNet-18)** through hardware-aware training.
-

7.0.4 4. Conclusion

DIANA proves that a hybrid digital-analog approach can overcome the individual limitations of both paradigms. By integrating a flexible digital core with a high-throughput analog core under a unified RISC-V controller, it enables efficient and accurate end-to-end inference for complete, complex neural networks at the edge.

8 L7 - Paper Summary: XpulpNN: Accelerating Quantized Neural Networks on RISC-V Processors Through ISA Extensions

8.1 Overview

This paper presents **XpulpNN**, a set of custom Instruction Set Architecture (ISA) extensions for RISC-V designed to accelerate low-bitwidth Quantized Neural Networks (QNNs). While existing microcontrollers (MCUs) typically support 8-bit or 16-bit precision, XpulpNN introduces native support for **sub-byte data types (2-bit and 4-bit)**. This effectively removes the performance bottleneck caused by manual data packing and unpacking in software, enabling high-efficiency Deep Learning inference on extreme-edge IoT devices.

8.2 Key Contributions

8.2.1 1. ISA Extensions (XpulpNN)

The authors extended the RV32IMCXpulpV2 ISA with instructions specifically targeting sub-byte SIMD (Single Instruction Multiple Data) operations:

- **Sub-byte SIMD Arithmetic:** Support for 4-bit (nibble) and 2-bit (crumb) operands, including addition, subtraction, average, and dot product.
- **Dot Product Unit:** A redesigned hardware unit that can compute the dot product of eight 4-bit or sixteen 2-bit elements in a single clock cycle.
- **Hardware Quantization Unit (pv.qnt):** A dedicated multicycle instruction that handles the “staircase” compression function (converting 16-bit intermediate results back to 2/4-bit activations). This unit replaces the computationally expensive software-based binary tree approach.

8.2.2 2. Microarchitectural Optimization

To maintain high performance without sacrificing energy efficiency:

- **Hardware Duplication vs. Sharing:** The design uses separate multiplier sets for different bitwidths to avoid the timing delays of shared logic, ensuring the critical path of the system remains unaltered.
 - **Power Management:** To mitigate the area overhead (11.1%), the designers implemented **clock gating** and **operand isolation** to disable inactive hardware units during general-purpose tasks.
-

8.3 Experimental Results

The researchers implemented the design in **22nm FDX technology** and compared the XpulpNN-extended core against a baseline RISC-V core and commercial ARM Cortex-M solutions.

8.3.1 Performance & Efficiency Gains

- **Speedup:** 4-bit kernels run **5.3x faster** and 2-bit kernels run **8.9x faster** than the baseline 8-bit SIMD processor.
- **Energy Efficiency:** The solution achieves a peak efficiency of **279 GMAC/s/W**, a **9x improvement** over the baseline.
- **Quantization Speed:** The dedicated `pv.qnt` instruction reduces the quantization phase overhead to just 4–11% of total execution cycles.

8.3.2 Comparison with State-of-the-Art

Metric	vs. ARM Cortex-M4/M7	vs. Baseline RISC-V
Performance	Up to 19.3x faster	Up to 8.9x faster
Energy Efficiency	Up to 354x better	Up to 9x better

8.4 Conclusion

The paper demonstrates that by integrating domain-specific ISA extensions into a programmable RISC-V core, it is possible to achieve **ASIC-like efficiency** while maintaining software flexibility. XpulpNN enables MCUs to scale performance linearly with lower quantization levels, providing a viable path for deploying complex Deep Learning models on resource-constrained IoT end-nodes.

9 L7 - Summary of ARM big.LITTLE Technology White Paper

This white paper details ARM’s **big.LITTLE technology**, a heterogeneous processing architecture designed to provide high performance for modern mobile devices while maintaining extreme energy efficiency.

9.0.1 Core Concept: Heterogeneous Coherency

The architecture utilizes two distinct types of processor clusters that share the same **Instruction Set Architecture (ISA)**, such as the ARMv7-A:

- **“LITTLE” Cores (e.g., Cortex-A7, A53):** Designed for maximum power efficiency with simple, in-order pipelines.
- **“big” Cores (e.g., Cortex-A15, A57):** Designed for maximum compute performance with complex, out-of-order, multi-issue pipelines.

The key to the system is hardware-level **Cache Coherency**, enabled by the **Cache Coherent Interconnect (CCI-400)**. This allows for the seamless and transparent transfer of data between clusters without relying on slow external memory transactions.

9.0.2 Software Execution Models

The paper identifies two primary methods for managing tasks in a big.LITTLE system:

1. **CPU Migration (Cluster Switching):** Pairs each big core with a LITTLE core. Only one core in the pair is active at a time based on load. This requires an equal number of cores in each cluster.
 2. **Global Task Scheduling (GTS):** A more sophisticated model (implemented as **big.LITTLE MP**) where the OS scheduler is aware of the different capacities of all cores. It can utilize all cores simultaneously, regardless of how many are in each cluster.
-

9.0.3 Dynamic Task Management in big.LITTLE MP

The **big.LITTLE MP** solution tracks the “load average” of individual threads in real-time to determine where they should run. It utilizes several migration techniques:

- **Fork Migration:** New threads default to a big core to ensure demanding tasks aren’t penalized at launch.
 - **Wake Migration:** Idle tasks becoming ready to run resume based on their last tracked load value.
 - **Forced Migration:** Long-running threads that exceed a specific “up migration threshold” are moved from LITTLE to big cores.
 - **Idle Pull Migration:** When a big core becomes idle, it “pulls” intensive tasks from LITTLE cores to maximize performance.
 - **Offload Migration:** Periodically migrates threads to LITTLE cores to utilize unused compute capacity when big cores are overloaded.
-

9.0.4 Performance and Efficiency Results

Benchmarks comparing a big.LITTLE system to a system using only big cores (Cortex-A15) show:

- **Power Savings:** Up to **75% CPU power savings** and **40% SoC power savings** across common use cases like audio, video, and gaming.
- **Performance Gains:** Significant improvements in threaded benchmarks (like CF-Bench and Antutu) by enabling the system to deploy more processors (e.g., 4 big + 4 LITTLE) than a standard quad-core big-only system.

10 L8 - Summary: A RISC-V Vector Processor with Tightly-Integrated Switched-Capacitor DC-DC Converters

10.1 Overview

This paper presents a high-efficiency RISC-V vector microprocessor implemented in **28nm FDSOI** technology. The primary innovation is the integration of **non-interleaved switched-capacitor DC-DC (SC-DCDC) converters** combined with an **adaptive clocking scheme**. This system allows the processor to operate reliably across a wide voltage range (0.5V to 1V) while maximizing energy efficiency for mobile applications.

10.2 Key Technical Features

10.2.1 1. RISC-V Processor Architecture

- **Core:** A 64-bit single-issue in-order scalar core implementing the open-source RISC-V ISA.
- **Accelerator:** A high-performance 64-bit vector accelerator supporting gather/scatter operations and fused multiply-add.
- **Memory:** Includes 16KB scalar instruction, 8KB vector instruction, and 32KB shared data caches. Custom **8T SRAM cells** were used to allow the memory to scale down to 0.45V reliably.

10.2.2 2. Integrated Power Management

- **SC-DCDC Converters:** Unlike traditional interleaved designs that minimize ripple, this work uses a **non-interleaved** design to eliminate charge-sharing losses, achieving higher peak efficiency.
- **Reconfigurable Topologies:** The converter uses 24 unit cells to generate four discrete average voltages (1.0V, 0.9V, 0.67V, and 0.5V) from 1.0V and 1.8V inputs.
- **Fast Transitions:** The system supports extremely fast voltage scaling transitions in under **20ns**.

10.2.3 3. Adaptive Clocking System

To handle the intentional voltage ripple created by the non-interleaved converters, the researchers implemented an **adaptive clock generator**.

- It uses a **Tunable Replica Circuit (TRC)** to mimic the critical path delay at the instantaneous voltage level.
 - By adjusting the clock frequency cycle-by-cycle based on the supply ripple, the system operates at the maximum possible frequency for any given moment, significantly improving performance compared to a fixed-frequency clock.
-

10.3 Measured Performance and Results

The chip was verified by booting Linux and running complex applications like Python. Key metrics include:

Metric	Value
Technology	28nm FDSOI
Total Die Area	2.37 mm ² (16% area overhead for converters)
System Efficiency	80–86% (Conversion efficiency)
Energy Efficiency	26.2 DP GFLOPS/W
Operating Range	93 MHz to 961 MHz
Power Consumption	8 mW to 173 mW

10.4 Conclusion

The design demonstrates that tight integration between power delivery and the processor core—specifically using “rippling” supplies with adaptive clocking—can outperform traditional regulation techniques. By avoiding off-chip passives and reducing charge-sharing losses, the system provides a compelling solution for energy-constrained mobile devices requiring high-performance vector processing.

11 L9 - Summary: eIQ Neutron - Redefining Edge-AI Inference

11.1 Overview

This paper introduces **eIQ Neutron**, a novel Neural Processing Unit (NPU) architecture and its accompanying compiler designed to maximize AI inference efficiency on resource-constrained edge devices. The authors argue that traditional metrics like **peak TOPS** (tera operations per second) are poor indicators of real-world performance. Instead, the eIQ Neutron focuses on **Hardware-Software (HW/SW) co-design** to minimize data movement and maximize compute utilization.

11.2 Key Architectural Principles

The eIQ Neutron architecture is built on three first principles for advanced CMOS nodes (16nm and below):

1. **Interconnect-Centric Design:** Minimizing global bandwidth needs because wire energy costs far exceed logic costs in deep sub-micron processes.
2. **Memory Locality:** Utilizing deep pipelining to tolerate latency and prioritizing local SRAM (Tightly Coupled Memory) access over power-hungry off-chip DRAM.
3. **Instruction Overhead Reduction:** Using a data-driven **systolic array** to perform large arithmetic operations per programming step, reducing the energy wasted on instruction fetching and decoding.

11.3 Innovations

11.3.1 1. Hardware Architecture

- **Dot-Product Systolic Engine:** Employs parallel, pipelined units that share operands to reduce input bandwidth requirements.
- **Output-Stationary Flow:** Avoids external memory accesses for high-precision 32-bit accumulators.
- **Data Engine:** A programmable pre-fetcher that supports spatial and temporal data reuse through multi-dimensional address generation.

11.3.2 2. Compiler and Software Stack

- **Spatial Tiling (Line vs. Depth Parallelism):** The compiler dynamically chooses between “Depth Parallelism” (sharing activations across filters) and “Line Parallelism” (sharing parameters across output lines) to optimize for different layer shapes.
- **Constraint Programming (CP):** The compiler uses formal CP models to solve complex scheduling and memory allocation problems. This optimizes the “tiling-scheduling-placement” space to hide memory latency.
- **Layer Fusion:** Interleaves the execution of multiple layers to reduce the need to write intermediate data back to off-chip memory, significantly lowering the SRAM footprint.

11.4 Experimental Results

The system was benchmarked on standard computer vision models (e.g., MobileNet, ResNet, YOLOv8) and compared against industry-leading NPUs:

- **Vs. Similar Resources:** Achieved an average **1.8x speedup** (up to 4x peak) compared to embedded NPUs with equal TOPS and memory.
- **Vs. Larger Systems:** Outperformed NPUs with double the compute/memory resources by up to **3.3x**.
- **Vs. High-End iNPUs:** A 2 TOPS Neutron system outperformed an 11 TOPS integrated NPU by **1.25x** on average, demonstrating that intelligent co-design is more effective than raw scaling.
- **Efficiency:** The design consistently achieved the best **Latency-TOPS Product (LTP)**, indicating superior performance-per-cost.

11.5 Conclusion

The paper concludes that raw TOPS should be replaced by metrics reflecting actual utilization. The eIQ Neutron demonstrates that through tight coupling of a flexible NPU architecture and a CP-based compiler, it is possible to achieve best-in-class efficiency for demanding edge-AI workloads like image segmentation and object detection.

12 License

This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License (CC BY-NC-SA 4.0) for KU Leuven students.

You are free to:

- **Share** — copy and redistribute the material in any medium or format with people inside of KU Leuven or being granted access to the source material.
- **Adapt** — remix, transform, and build upon the material if you are a KU Leuven student.

Under the following terms:

- **Attribution** — You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.
- **NonCommercial** — You may not use the material for commercial purposes.
- **ShareAlike** — If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.

For the full legal text of the license, please visit: <https://creativecommons.org/licenses/by-nc-sa/4.0/legalcode>

12.1 Modification

To contribute to this work or any other one from this project please find more information at the Github repository.

12.2 Take down

If you are a professor or representing any official party and wish to take down this work, you can submit a takedown request at this url: https://github.com/Tfloow/ESATSummary/issues/new?template=takedown_notice.yml

Some Rights Reserved 2026 Authors of the Summary, Professors of the Course and possible book's authors. Some Rights Reserved.