
Image Analysis

An Open-Source Summary

Thomas Debelle



April 29, 2026

Contents

1	Classic Image Analysis	2
2	Introduction	2
2.1	The basics	2
2.1.1	Light	2
2.1.2	Perception	2
2.1.3	Interaction with matter	2
3	Acquisition of images	6
3.1	Image formation - Lens	7
3.2	Sensor array	8
3.2.1	CCD vs CMOS	8
3.2.2	Colour cameras	9
3.3	Perspective Projection - Image Plane	10
3.3.1	Projections limitations	11
3.4	Deviations from the lens model	13
3.4.1	Radial distortion	14
3.4.2	Chromatic aberration	15
3.4.3	Photometric camera model	15
4	Sampling and Quantisation	15
4.1	Sampling	17
4.1.1	Integrate brightness over the cell window	17
4.2	Aliasing	20
4.2.1	Binomial filters	23
4.2.2	Median	24
4.2.3	Anisotropic filters - all-in-on of filtering	25
4.3	Enhancement and Feature Detection	26
4.3.1	Edges	26
5	Modern Image Analysis	26
6	License	27
6.1	Modification	27
6.2	Take down	27

1 Classic Image Analysis

2 Introduction

The key take away is that vision is subjective to the human being and themselves. It is the results of continuous evolution and survival of the fittest. We are more sensitive to light and certain colors, to motion, ... And this fact will be used for efficient image representation and analysis.

We can think of the classic illusion where our brain will interpret certain lengths of an arrow longer than others. Context also does matter, even in a blurry picture our brain will train to make up for what he sees.

2.1 The basics

We first need light to have vision. Light is influenced by object and create perception. ADD image slide 51. The light source bounce on the object and some scattered ray will get trough the 2D Image plane, which will be casted on the sensor array thanks to the lens system.

2.1.1 Light

We will look at it from a physical optic (not geometrical (constant speed and straight path) nor quantum). Light is an Electro-magnetic wave which we assume it to be planar. aka $\frac{|E|}{|H|} = \eta$ in simple term, light is perpendicular. This is a **self-sustaining process** that is characterized by:

1. λ : wavelength
2. Direction
3. E : Amplitude
4. φ : phase
5. Direction of polarisation

Violet: low frequency ($380nm$); Red: high frequency ($760nm$). Remember $\lambda = \frac{c}{f}$.

2.1.2 Perception

We don't perceive light equally. Humans have 3 cones that are not evenly spread apart. But the human vision is trained to compensate for it so everything feels similar in intensity to us. Some animals can have more cones.

2.1.3 Interaction with matter

Table 1: The four types of interaction

Phenomenon	Example
Absorption	Blue water, green leaf
Scattering	blue sky, red sunset

Phenomenon	Example
Reflection	colored in k
Refraction	dispersion by a prism

2.1.3.1 Absorption Dissipation of λ specific for the medium. Resonance in molecules.

2.1.3.2 Scattering When light (radiation/particles) deviates from a straight trajectory by local non-uniformities. Type depends on the relative size of particles and wavelengths:

1. **Rayleigh:** small (λ dependent)
2. **Mie:** comparable (weakly λ dependent)
3. **Non-selective:** large (λ independent)

We ignored transmission and diffraction.




Figure 1: Wavelength dependence for scattering

We need scatter to get the bluesky (Rayleigh with Tyndall effect) or it would be dark. Coloured cloud from volcanic eruption = Mie (particle all over the air). Grey cloud is due to non-selective. Less scatter for lower frequency light (infrared).




Figure 2: Mirror reflection

2.1.3.3 Reflection With reflection, the notion of polarization is important. The polarization effect creates the fact that a reflected wave on certain surfaces can be reflected with all but one direction of the electric field. This effect is more prominent on **non-metallic** surfaces!




Figure 3: Non-metallic reflection - strong reflectors




Figure 4: Metallic reflection - Red dot: Brewster angle (polarizer) - Blue zone: Grazing angles

Rough surfaces will lead to **diffuse** reflection as the reflection angles will be all over the place. On the other hand, smooth surfaces give **specular** reflection. Both **diffuse** and **specular** can also happen all together.

Reflectance varies through λ and different side of vegetation can typically give various reflection¹

2.1.3.4 Refraction Typically, refraction which is the bending of the light and dispersion which is a frequency dependency.




Figure 5: Refraction - governed by Snell's law property of the medium and material

¹This fact can be used for spectral reflectance of vegetation since it forms a signature

3 Acquisition of images




Figure 6: Image acquisition model

Controlling the lights is essential for acquisition as it is the main vector to acquire images. Several techniques exist

Type	Description	Examples
Back-lighting	Light source behind object (better with a diffuser plate)	High contrast silhouette. Binary vision , inspection
Directional	Tilt lights to generate sharp shadow of specular reflection (rough surfaces)	Detection of cracks
Diffuse	Uniform lighting, all directions contribute to the diffusion reflection but specular spike due to spike	Detection of cracks
Polarized: Contrast diffuse and specular	diffuse makes light non-polarized while specular keeps polarized. Put polarizer/analyzer orthogonal, better for glare (high dynamic range) issue.	Detection of cracks
Polarized: Contrast dielectrics and metal	Dielectric at Brewster angle can be used as a polarizer which is not possible for metallic surface (see previous chapter).	So use a polarizer to suppress specular reflection from dielectric OR distinguish between specular and dielectric surfaces
Coloured	Highlight region of similar colour. Use laser (monochromatic light), differentiation between specular and diffuse. Comparing colours. Spectral sensitivity of a sensor.	
Structured	Spatially or temporally modulated light pattern. Projection on a 3D object will distort the projection pattern	3D reconstruction
Stroboscopic	High intensity light flash. Eliminate motion blur.	For high speed inspection

Note about the dark and bright field. Mostly relevant for the Directional/Diffuse lighting. *Dark field*, when the camera is placed outside the specular reflection area of the normal area, only specular reflection (different orientation from normal)

will show up. On the other hand, *Bright field*, we place camera inside this specular reflection light area and the non-normal area will appear dark.

3.1 Image formation - Lens

We use the pinhole model in this class:




Figure 7: Pinhole model

From this we can use the thin-lens equation which assumes:

- Spherical lens surfaces
- Incoming light $\pm \parallel$ to axis
- Thickness $\ll$ radii
- Same refractive index on both sides





Figure 8: Thin-lens equation



$$\Delta Z_0^- = Z_0 - Z_0^- = \frac{Z_0(Z_0 - f)}{Z_0 + f \frac{d}{b} - f}$$

Figure 9: Depth of field of the thin lens

There is one focal point. Strike a balance between incoming light (d) and large depth-of-field. Smaller depth-of-field is present in microscopes for example.

3.2 Sensor array

3.2.1 CCD vs CMOS

We have 2 types of sensors:

- CCD: Charge-coupled device
 - Old school
 - Charges must hop from tile to tile
 - Cover the total pixel (high fill rate)
 - Expensive
 - Blooming (due to the charge hopping)
 - High power consumption
- CMOS: integrated
 - Cheap, Can put other component on the chip (integration)
 - Read line by line (like RAM = random access)
 - Smaller
 - Low power
 - Same sensor as CCD but each has their own amplifier = more noise + not covering the full space = lower fill rate / sensitivity

They both work in 2 steps:

1. Photon to electron (charge) conversion
2. Electron (charge) to voltage conversion




Figure 10: CCD vs CMOS

3.2.2 Colour cameras

3.2.2.1 1. Prism The idea is to split the light using *dichroic prism* to split it in 3 colors. This requires really good precision but allows for good color separation.

3.2.2.2 2. Filter mosaic Back to the *bayer pattern* introduced earlier. Good for tricking the human eyes. Each pixel will have a filter in front to follow the bayer pattern. This is good but can create some aliasing especially when an edge is right on a pixel array causing some extra inexistant colors (aliasing).




Figure 11: Aliasing with bayer pattern

We loose actual resolution since now we need 2x2 pixel array to get 1 colored pixel ! We have to stack even more density but need to capture more lights → Use microlenses.

3.2.2.3 3. Filter wheel Here instead of having 3 filters like in a filter mosaic, we have a full wheel with many filters. Behind the wheel, there is a static sensor (**only static scenes**).

3.2.2.4 Summary

Table 3: Summary of the 3 coloured camera

Type	Prism	Mosaic	Wheel
# Sensors	3	1	1

Type	Prism	Mosaic	Wheel
Resolution	High	Medium	Good
Cost	High	Low	Medium
Framerate	High	High	Low
Artefacts	Low	Aliasing	Aliasing (just motion of wheel)
Bands (color)	3	3	3 or +
Application	High-end cameras	Low-end cameras	Scientific applications

3.2.2.5 Improvements One issue is the Aliasing problem for object in motion. Since we read line by line in CMOS camera, for fast turning object (turning at $\geq f_s/2$) we have aliasing. To solve this, we need a global shutter instead of reading line by line.

3.3 Perspective Projection - Image Plane

Here, we revisit the pinhole model where now, we first project everything on an object plane and then this will get through the lens. The center of the lens is the center of the projection, we have an virtual image plane which is also called **perspective projection** (we flattened out everything from the outside on it):




Figure 12: Camera projection model




Figure 13: 3D view of the camera projection model

The origin is the center of the projection (X_c row-pixel, Y_c column-pixel space), the Z_c axis correspond with the optical axis.

If we look at the u, v space we have the following relationship, which comes from the “similar triangles property” where you look once from above (for X case) and one from the side (for Y case).

$$u = f \frac{X_c}{Z_c} \quad v = f \frac{Y_c}{Z_c}$$

We can then say that if Z is constant, we can write $x = kX_c$ and $y = kY_c$ where $k = f/Z$. Orthographic projection + a scaling. This forms a good approximation if objects are small compared to their distance from the camera.

3.3.1 Projections limitations



Figure 14: Pseudo-orthographic vs Perspective projection

The perspective projection model is incomplete. What about world coordinate ? What if we want our u, v coordinates as a row and column numbers (like in a matrix of pixels).

This part of the class that does not form exam material but helps with understanding where does those matrices come from.

3.3.1.1 The (u, v) projection to pixel We can see a matrix of pixel instead of the (u, v) projection. We just need 3 informations: 1. Center of matrix x_0 and y_0 , 2. Width of array k_x and 3. Height of array k_y giving us:

$$\begin{cases} x &= k_x u + x_0 + sv \\ y &= k_y v + y_0 \end{cases}$$

Note: the s is for skew, we have the aspect ratio k_y/k_x ²

3.3.1.2 The global world coordinate problem Those equations are based on the previous $u = fX/Z$ and $v = fY/Z$. Here $\langle, \rangle$ represents the dot product $\cdot$ between 2 vectors.

²I thought it was the opposite...




Figure 15: Projection matrices

The $P - c$ represents the dotted vector and C being the camera position in a global coordinate system (the camera is no longer stuck to the 0,0 point). We then use the dot product between one of its axes and the vector $P - C$ to obtain the value in this specific vector dimension (a bit like the X before). Repeat it for Y and Z .

3.3.1.3 Homogeneous coordinates Capturing a 2D image requires to navigate a 3D world (and so on for 3D and 4D). This is important in computer vision where we have to deal with projection where we **have to** divide a point by z !

$$\tau \begin{pmatrix} x \\ y \\ z \end{pmatrix} \rightarrow_{\text{capturing/projection}} \begin{pmatrix} x/z \\ y/z \end{pmatrix}$$

3.3.1.4 Putting it all together We can rewrite the previous big dot/inner product using the homogeneous coordinate (removes division) as:

$$\tau \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} fr_{11} & fr_{12} & fr_{13} \\ fr_{21} & fr_{22} & fr_{23} \\ r_{31} & r_{12} & r_{13} \end{pmatrix} \begin{pmatrix} X - C_1 \\ Y - C_2 \\ Z - C_3 \end{pmatrix}$$

we can then apply the pixel coordinate system on this equation:

$$\tau \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = \begin{pmatrix} k_x & s & x_0 \\ 0 & k_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} fr_{11} & fr_{12} & fr_{13} \\ fr_{21} & fr_{22} & fr_{23} \\ r_{31} & r_{12} & r_{13} \end{pmatrix} \begin{pmatrix} X - C_1 \\ Y - C_2 \\ Z - C_3 \end{pmatrix}$$

Back to important material for exam.

we can further simplify with

$$\tau \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = \overbrace{\begin{pmatrix} k_x & s & x_0 \\ 0 & k_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{pmatrix}}^{K \text{ Calibration matrix}} \overbrace{\begin{pmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{12} & r_{13} \end{pmatrix}}^{R \text{ Camera rotation}} \begin{pmatrix} X - C_1 \\ Y - C_2 \\ Z - C_3 \end{pmatrix}$$

We then define the following vectors:

$$p = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} \quad P = \begin{pmatrix} X \\ Y \\ Z \end{pmatrix} \quad \tilde{P} = \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} \quad (1)$$

Which then yields the simplified version:

$$\rho p = K R^\top (P - C) \quad \text{some non-zero } \rho \in \mathbb{R}$$

$$\text{or, } \rho p = K(R^\top | - R^\top C) \tilde{P}$$

$$\text{or, or, } \rho p = (M|t) \tilde{P} \quad \text{with rank } M = 3$$

3.4 Deviations from the lens model

We assume:

1. All rays from a point are focused onto 1 image point
2. All image points in a single plane
3. Magnification is constant

If those are not respected, we have **aberrations**! Aberrations can be **geometrical** (small for paraxial rays) or **chromatic** (refractive index function of λ - Snell's law)




Figure 16: Geometrical aberrations

They all cause issues on the focus point of the image plane. Spherical appears when the bending of light is not uniform throughout the lens, typically outer ray converges sooner (small focal length) than inner one. They do not converge all together. Coma is a bit like spherical but this time with a tilted ray. Astigmatism is problems between the *tangential* and *sagittal* (90° rotation from tangential) focus point which are not equal.

3.4.1 Radial distortion

By far the most important type of distortion which causes “fish-eye” look or pincushion. It creates various *magnification* for different angles of inclination. The pixel either sinks towards the center or slides outside along the line connection the center and the original pixel.

magnification different for different angles of inclination



Figure 17: Radial Distortion

We can model this sliding distance d with the following equation and thus can be corrected. Some algorithm looks at human made structure (straight lines) to establish the effect.

$$d = (1 + \kappa_1 r^2 + \kappa_2 r^4 + \dots)$$

3.4.2 Chromatic aberration

Rays of different λ will focus in different planes. This creates various colours especially along sharp edges. This **can't** be removed completely but *achromatization* can be achieved by selecting the right type of glasses.




Figure 18: Achromatization

3.4.3 Photometric camera model

We want to convert the object radiance into a pixel grey level. This is a 2-step process with our idea of the image plane. We first project the radiance of the object to the image, then this image irradiance projects to a pixel in a grey level. The **cos⁴** law governs the radiance. The more an object is off-axis from the center by an angle θ the more it reduces its irradiance in the pixel. This materializes with **natural vignetting**³ (like the old instagram filter).

$$I = R \frac{A_l}{f^2} \cos^4 \theta$$

This will get converted in a grey pixel with:

$$pix = gI^\gamma + d$$

With the g of the camera, d the dark reference of the camera. The gamma γ is a non-linear relationship which changes the mid-tone without really changing the pure black or white. A value lower than 1 makes the image lighter and vice-versa. Nowadays, cameras have become quite linear meaning that $\gamma = 1$.

4 Sampling and Quantisation

We live in an analog world but our computers and digital world is finite. So we have to sample (Discretize points) and quantize (attribute them a finite value).

From what we have seen and learned before, we can trick the human eye to see a higher quality image than it is actually. For example the display on a phone has a RGB pixel ratio as 2:1:1 because the human is less sensible to green. We call

³optical vignetting appears due to an obstruction of the lens by an element

this grid the **bayer filter**. Other representation exist which can be more exotic, e.g.: hexagonal (more isotropic, similar structure as in the retina, no connectivity ambiguities).

We have quantization where we usually talk about levels. Typically a n bits representation has 2^n levels. We can also apply fancier quantization scheme instead of simple linear one. The most classic representation is the RGB888 (1 byte per pixel for each color).





Figure 19: Various quantization scheme

4.0.0.1 HDR High Dynamic Range (HDR) is a technology where we have various luminance on the screen giving really good contrast in dark and bright area. Usually, we combine multiple photos with various gains and then recombine it all together with the various exposure.

4.1 Sampling

4.1.1 Integrate brightness over the cell window



$$o(x', y') = \iint i(x, y) p(x - x', y - y') dx dy$$

This is **convolution**: $i(x, y) * p(-x, -y)$

Figure 20: Integrating over a pixel cell

This forms a convolution. Convolution are **linear**, $f_1 \rightarrow g_1 \quad f_2 \rightarrow g_2 \Rightarrow af_1 + bf_2 \rightarrow ag_1 + bg_2$, and **shift-invariant**, $f(x, y) \rightarrow g(x, y) \Rightarrow f(x - a, y - b) \rightarrow g(x - a, y - b)$.

Convolutions are also commutative and associative as proven in the Fourier domain.

$$\mathfrak{F}(u) = \int_{-\infty}^{\infty} f(t) e^{-2i\pi xu} dx$$

For the 2D case, we have $e^{-2i\pi(ux+vy)}$ which has for Euler form $\cos(2\pi(ux + vy)) + i \sin(2\pi(ux + vy))$. Now, the u and v represents the frequency along the x and y dimensions. To obtain the λ of the signal we need to take the norm of the frequency for both axis inverted:

$$\lambda = \frac{1}{\sqrt{u^2 + v^2}}$$

As in 1D, we can do Fourier transform to decomposition the image in a bunch of frequencies. As in 1D, we have a complex results where the magnitude is $|\mathfrak{F}(u, v)| = \sqrt{\mathbb{R}(\mathfrak{F}(u, v))^2 + \mathbb{I}(\mathfrak{F}(u, v))^2}$ and the phase angle $\angle \mathfrak{F}(u, v) = \arctan(\mathbb{I}(\mathfrak{F}(u, v))/\mathbb{R}(\mathfrak{F}(u, v)))$.

$$\mathfrak{F}(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-2i\pi(xu+yv)} dx dy$$




Figure 21: Fourier decomposition of images




Figure 22: Repetitive patterns in an image creates repetitions in frequency domain

Removing patterns peak in frequency allows to remove periodic background.





Figure 23: Mixmatch of phase and magnitude

4.1.1.1 Properties of spatial-frequency domain

Spatial domain	Frequency domain
Real	Real part → even; Imaginary part → odd
Real, even	Real, even (cosine)
Real, odd	Imaginary, odd (sine)

Important to remember that a convolution in spatial domain is a multiplication in the frequency domain and vice versa.



We have the base image (the dot) and the the transfer function, MTF R, which contains the amplification and phase shift information to every frequency in the input image I.

4.1.1.2 Integrating with the convolution Re-using the integration over a cell window:

$$o(x', y') = \int \int i(x, y) p(x - x', y - y') dx dy \leftrightarrow i(x, y) * p(-x, -y)$$

We can then transform this $p(x, y)$ round-point (as seen in the previous figure) in a $P(u, v)$

$$P(u, v) = \int_{-\infty}^{\infty} e^{-i2\pi(ux+vy)} p(x, y) dx dy \quad (2)$$

$$= wh \left(\frac{\sin(\pi w u)}{\pi w u} \right) \underbrace{\left(\frac{\sin(\pi h v)}{\pi h v} \right)}_{\text{sinc}} \quad (3)$$

The sinc function is predominantly low-pass **and** non-causal (no phase shift) and it has phase reversal.

4.2 Aliasing


This corresponds to a 2D Dirac train convolution on the grid like:

$$f(a, b) = \int \int f(x, y) \delta(x - a, y - b) dx dy$$

The multiplication with the 2D pulse train with a spacing of distance w and h :

$$\sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} \delta(x - kw, y - lh) \quad (4)$$

$$\longleftrightarrow \frac{1}{wh} \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} \delta(x - k\frac{1}{w}, y - l\frac{1}{h}) \quad (5)$$



Sampling in one domain leads to a periodic repetition in the other, here space $\rightarrow$ frequency

Figure 24: Basic sampling theorem

Also, we need to do a Discrete Fourier Transform, DFT, to discretize the space and frequency. We can also use FFT to bring the complexity of the algorithm from N^2 to $N \log_2 N$.

$$\sum_{m=0}^{M-1} \sum_{n=0}^{N-1} F(k, l) e^{2\pi i \left(\frac{km}{M} + \frac{ln}{N} \right)} \quad (6)$$

$$\longleftrightarrow F(k, l) = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f(m, n) e^{-2\pi i \left(\frac{km}{M} + \frac{ln}{N} \right)} \quad (7)$$




Figure 25: Periodicity assumed in both domains, might introduce false high frequencies at image boundaries




Figure 26: We can have false high frequencies due to sampling




Figure 27: Solution on periodicity

Again, the Shannon theorem is important in 2D as well such that the sampling distance w, h along the x, y axis must follow:

$$w \leq \frac{1}{2u_b} \quad h \leq \frac{1}{2v_b}$$

ABOVE SHOULD BE THE START OF CHAPTER ABOUT ENHANCEMENT AND EDGES

P.41

4.2.0.1 Convolution filters It is based on the insight that noise affects high frequencies the most due to their low SNR. To build a low-pass filter, we will use an averaging kernel such:

$$\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix} = \underbrace{\begin{pmatrix} 1 & 1 & 1 \end{pmatrix} * \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}}_{\text{separable filter}}$$

The two $(1, 1, 1)$ vectors can be seen in the *frequency* domain as $1 + 2 \cos(2\pi u)$ for the horizontal domain or $1 + 2 \cos(2\pi v)$

in the vertical one. The intuition for this is to look at the vector as a 1D image centered around the middle value. Finally, the convolution can be simplified in the frequency domain as:

$$(1 + 2 \cos(2\pi u))(1 + 2 \cos(2\pi v))$$




Figure 28: Representation in 3D of the function

The function is purely real and low-pass with some phase reversal! We can again derivate something similar for a 5x5 kernel:

$$(1 + 2 \cos(2\pi u) + 2 \cos(4\pi u))(1 + 2 \cos(2\pi v) + 2 \cos(4\pi v))$$

With this one, the phase reversal is even more significant, the higher the frequency the more it looks like a sinc function. This creates more ripples and disturbing effects. Hard to get no ripple and a real low-pass for higher order functions.

From all those experiments we can draw 3 conclusions:

1. Separable filters can be implement more efficiently
2. Better to go in the frequency domain for large kernel
3. Keep it simple, mask of boolean or power of 2 are cheap!

4.2.1 Binomial filters

They are here to solve the ripple and low-pass issue of the convolution filters presented before. This guarantees no phase reversal.

$$\begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 2 & 1 \end{pmatrix} * \begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix} \longleftrightarrow (2 + 2 \cos(2\pi u))(2 + 2 \cos(2\pi v))$$




Figure 29: Gaussian filters are limit case of binomial filters

Note that noise removal linear filters will exhibit blurring due the fact they cannot difference noise and actual information.

4.2.2 Median

Instead of averaging over using a kernel of ones, we take the median - the middle point of neighborhood of pixel. We don't create a new grey level, we take the one in the middle.




Figure 30: Median filters principle - credits: southampton.ac.uk

One advantage of a median filter is the fact it is robust against outliers and it will preserve the contrast instead of smoothing out all the values and making the image more blurry.

- Discard outliers
- Preserve discontinuities

But, it will lose in details (especially small areas). Moreover, give a patchy effect instead of a smooth uniform color.

4.2.3 Anisotropic filters - all-in-one of filtering

This filter does:

1. Binomial smoothing between edges
2. Unsharp masking at edges

For this, we need to cover edge detection

Table 5: Summary table of the enhancement procedures

Operation name	Description	Linear ?
Histogram equalisation	Re-sort histogram to create a more linear Cumulative Probability Density	V
Unsharp mask		V
Inverse Filtering	Realize the opposite operation. Need to carefully handle low SNR with a constant.	V
Low pass	Removes noise since it affects more high-frequencies due to low SNR	V
Binomial	Like low-pass but avoids phase reversal	V
Median	Take the middle point instead of the mean	X
Anisotropic	Binomial between edges - Unsharp at edge	X

4.3 Enhancement and Feature Detection

4.3.1 Edges

Definition:

Edge correspond to image locations with strong intensity changes due to:

1. Reflectance
2. Orientation of shape normals (its surface) - the object's edge
3. Illuminations (shadow)

5 Modern Image Analysis

6 License

This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License (CC BY-NC-SA 4.0) for KU Leuven students.

You are free to:

- **Share** — copy and redistribute the material in any medium or format with people inside of KU Leuven or being granted access to the source material.
- **Adapt** — remix, transform, and build upon the material if you are a KU Leuven student.

Under the following terms:

- **Attribution** — You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.
- **NonCommercial** — You may not use the material for commercial purposes.
- **ShareAlike** — If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.

For the full legal text of the license, please visit: <https://creativecommons.org/licenses/by-nc-sa/4.0/legalcode>

6.1 Modification

To contribute to this work or any other one from this project please find more information at the Github repository.

6.2 Take down

If you are a professor or representing any official party and wish to take down this work, you can submit a takedown request at this url: https://github.com/Tfloow/ESATSummary/issues/new?template=takedown_notice.yml

Some Rights Reserved 2026 Authors of the Summary, Professors of the Course and possible book's authors. Some Rights Reserved.